

Resources Optimization in Kubernetes

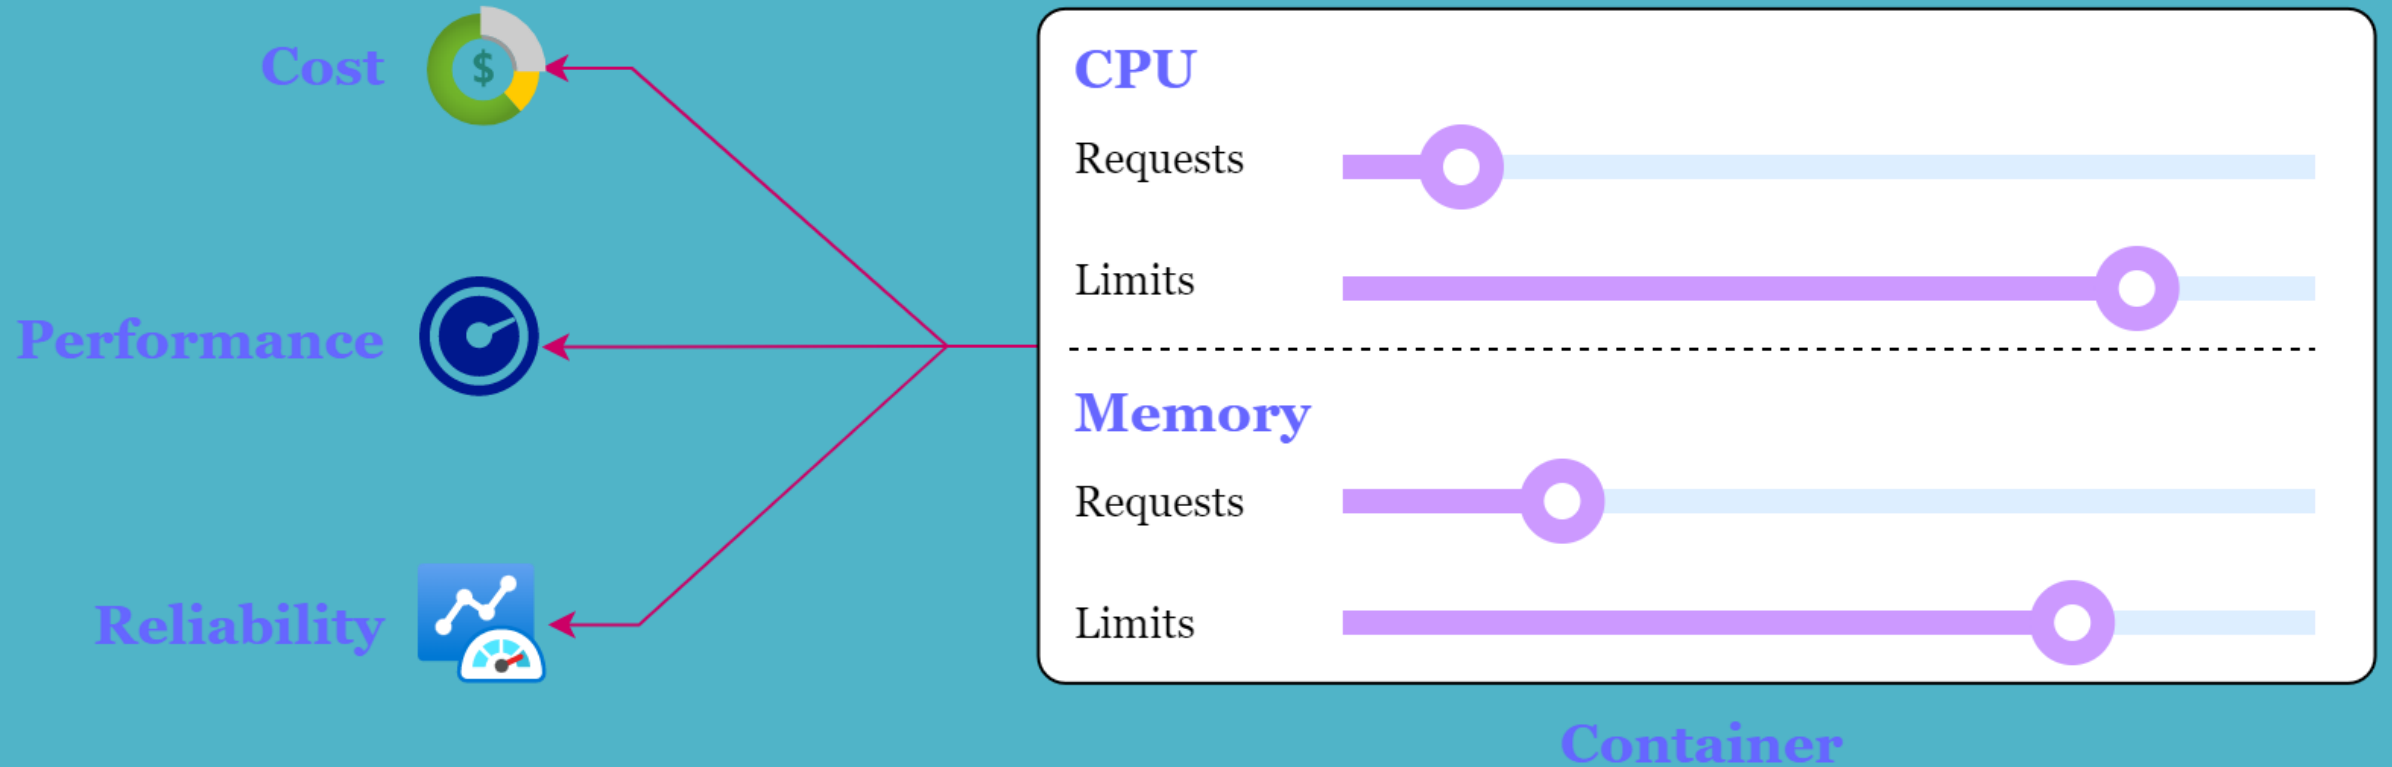
Trịnh Quốc Việt

VietOpenInfra Meetup 33rd – 8/2023

Content

- Requests & Limits
- Monitor resources
- How to optimize your resource
- Tools & demo

Limits & Requests



Limits & Requests

Application
types

Balanced

CPU



Memory



CPU Intensive

CPU




Memory



Memory Intensive

CPU



Memory



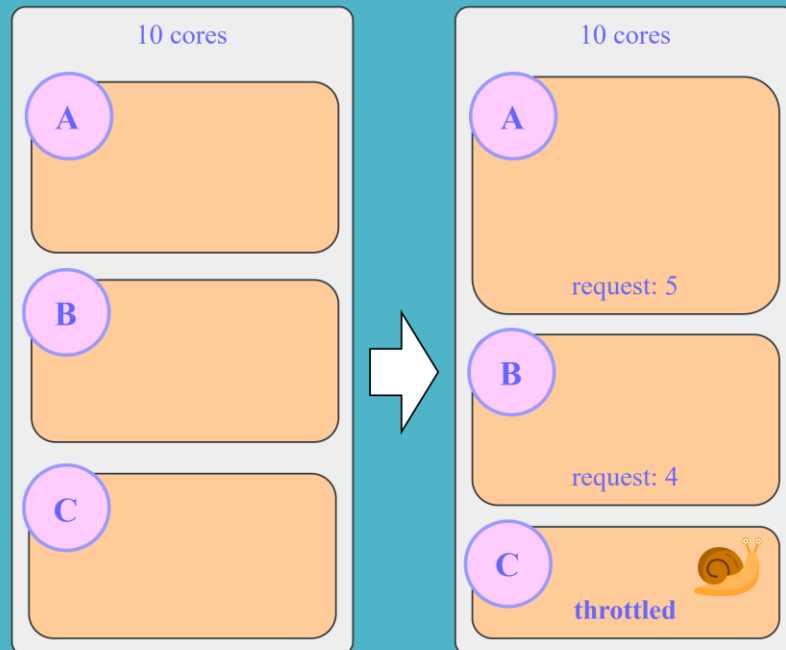
Limits & Requests

CPU

computing processing time

compressible resource

throttling

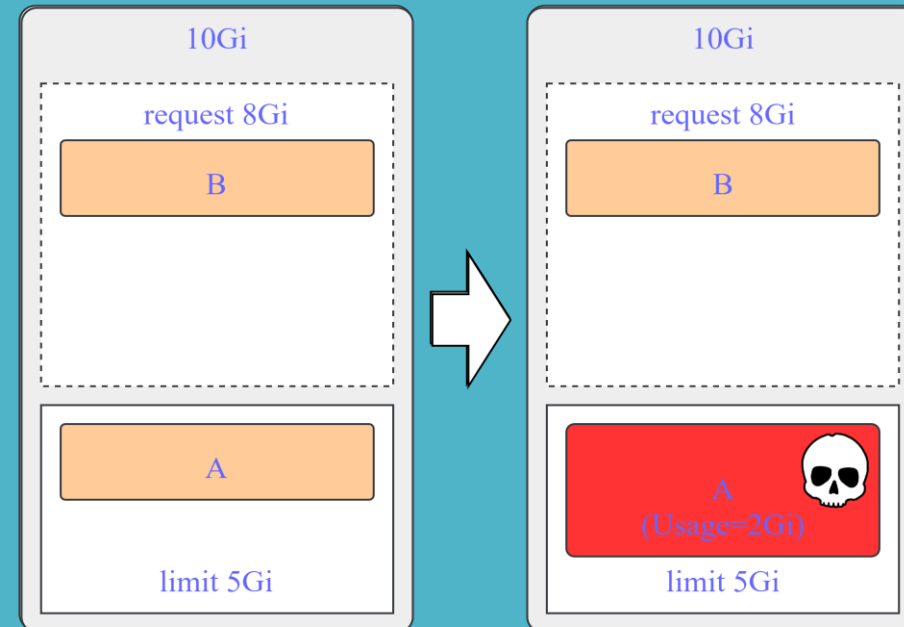


Memory

Measure in Mi, Gi..

Non-compressible resource


Out Of Memory killed




Limits & Requests

Requests

CPU




Too low
(Throttling)




Ideal value
(requests ~ usage)

Too high
(resource wasted)

Memory



Too low
(leak of memory)




Limits & Requests

Limits

CPU




Too low
(Throttling)




Ideal value
(requests ~ usage)

Too high
(consumes all
node's resource)


Memory



Too low
(OOM killed)




Limits & Requests - Scheduling




Limits & Requests - QoS


BestEffort



Burstable



Guaranteed



Priority

Monitor resources

What we need to monitor

CPU throttling

OOM Kill

Wasted resources

Actual usages

Monitor resources

Wasted resources

cluster/nodes level

Namespace level

Application level

Container level

Monitor resources

Wasted resources - node level

Cluster	Name		Role	Instance Type	Version	CC	MC	CPU	Memory (GiB)
cluster	staging	01	worker	rke2	v1.24.10+rke2r1	30	47.1 GiB	<div><div>8.7</div><div>19.3</div><div>28.5</div></div>	<div><div>32.1</div><div>28.1</div><div>44.1</div></div>
cluster	staging-	02	worker	rke2	v1.27.1+rke2r1	30	47.1 GiB	<div><div>4.8</div><div>19.6</div><div>28.5</div></div>	<div><div>21.5</div><div>23.0</div><div>44.1</div></div>
cluster	staging-	03	worker	rke2	v1.24.10+rke2r1	30	47.1 GiB	<div><div>11.7</div><div>18.4</div><div>28.5</div></div>	<div><div>29.8</div><div>29.3</div><div>44.1</div></div>
cluster	staging-	04	worker	rke2	v1.24.10+rke2r1	30	47.1 GiB	<div><div>9.8</div><div>21.5</div><div>28.5</div></div>	<div><div>26.9</div><div>32.6</div><div>44.1</div></div>
cluster	staging	05	worker	rke2	v1.24.10+rke2r1	30	78.6 GiB	<div><div>7.6</div><div>18.3</div><div>28.5</div></div>	<div><div>30.3</div><div>17.5</div><div>75.6</div></div>
cluster	staging	01	worker	rke2	v1.24.10+rke2r1	12	23.5 GiB	<div><div>5.0</div><div>2.6</div><div>10.5</div></div>	<div><div>16.1</div><div>2.3</div><div>20.5</div></div>
cluster	staging-	02	worker	rke2	v1.24.10+rke2r1	12	23.5 GiB	<div><div>3.8</div><div>2.6</div><div>10.5</div></div>	<div><div>14.6</div><div>2.3</div><div>20.5</div></div>
cluster	staging	03	worker	rke2	v1.24.10+rke2r1	12	23.5 GiB	<div><div>3.8</div><div>2.5</div><div>10.5</div></div>	<div><div>13.7</div><div>2.2</div><div>20.5</div></div>

Monitor resources

Wasted resources - namespace level

All Namespaces

ID	Age	Cluster	P	CR	MR	CPU	Memory (MiB)	Cost	Slack Cost
-staging	✓	cluster	44	16.9	31.6 GiB	<div><div>1.3</div><div>16.9</div></div>	<div><div>7,185</div><div>32,400</div></div>	3.86	2.97
-staging	✓	cluster	58	14.8	35.4 GiB	<div><div>0.55</div><div>14.8</div></div>	<div><div>17,474</div><div>36,300</div></div>	4.16	2.25
longhorn-system	✓	cluster	72	13.15	0 Bytes	<div><div>27.95</div><div>13.15</div></div>	<div><div>19,707</div><div>0</div></div>	2.27	0.00
istio-system	✓	cluster	21	12.52	7.4 GiB	<div><div>0.29</div><div>12.52</div></div>	<div><div>3,864</div><div>7,564</div></div>	2.16	1.56
-staging	✓	cluster	18	10.9	16.2 GiB	<div><div>0.29</div><div>10.9</div></div>	<div><div>3,831</div><div>16,640</div></div>	2.01	1.54
kube-system	✓	cluster	44	6.975	7.3 GiB	<div><div>5.2</div><div>6.98</div></div>	<div><div>33,128</div><div>7,480</div></div>	1.40	0.47
-staging	✓	cluster	4	3.58	7.6 GiB	<div><div>0.05</div><div>3.58</div></div>	<div><div>1,191</div><div>7,826</div></div>	0.81	0.33
argocd	✓	cluster	10	3.4	3.4 GiB	<div><div>4.06</div><div>3.4</div></div>	<div><div>2,812</div><div>3,512</div></div>	0.59	0.27


Monitor resources

Wasted resources - Application level

ID	Team	AP	C	P	CR	MR	CPU	Memory (MiB)
			1	117	20.656	8.0 GiB	<div><div>25.0520.66</div></div>	<div><div>37,7508,220</div></div>
istiod			1	10	10.0	4.9 GiB	<div><div>0.110.0</div></div>	<div><div>1,7015,000</div></div>
cache-service			1	2	4.0	3.9 GiB	<div><div>0.014.0</div></div>	<div><div>4304,000</div></div>
autoservice			1	4	4.0	3.9 GiB	<div><div>0.094.0</div></div>	<div><div>9534,000</div></div>
lis-service			1	1	3.0	2.9 GiB	<div><div>0.03.0</div></div>	<div><div>803,000</div></div>
nifi			1	1	2.53	6.5 GiB	<div><div>0.012.53</div></div>	<div><div>876,674</div></div>
up-service			1	1	2.5	3.4 GiB	<div><div>0.012.5</div></div>	<div><div>1263,500</div></div>
userservice			1	4	2.4	3.1 GiB	<div><div>1.282.4</div></div>	<div><div>1,8713,200</div></div>
proj-service			1	1	2.1	1.5 GiB	<div><div>0.012.1</div></div>	<div><div>2031,500</div></div>

Monitor resources

Container average CPU usages



How to optimize your resource

*Always set requests
& limits*

*Set the right requests
& limits*

*Understand well
your app*

*Monitor resource
periodically*

Tools & demo

Metrics server

kube-capacity

kube-resource-report

prometheus/grafana

Q & A

Ask me any questions

Thank you!